


Computer Vision Assignment

Hugo Collins

40446925



Dataset Preparation

```
Folder: n02808440 - Label: 0 - Class Name: bathtub, bathing tub, bath, tub
Folder: n01742172 - Label: 1 - Class Name: boa constrictor, Constrictor constrictor
Folder: n01443537 - Label: 2 - Class Name: goldfish, Carassius auratus
Folder: n02909870 - Label: 3 - Class Name: bucket, pail
Folder: n02481823 - Label: 4 - Class Name: chimpanzee, chimp, Pan troglodytes
Folder: n02364673 - Label: 5 - Class Name: guinea pig, Cavia cobaya
Folder: n02788148 - Label: 6 - Class Name: bannister, banister, balustrade, balusters, handrail
Folder: n01984695 - Label: 7 - Class Name: spiny lobster, langouste, rock lobster, crawfish, crayfish, sea crawfish
Folder: n02504458 - Label: 8 - Class Name: African elephant, Loxodonta africana
Folder: n03126707 - Label: 9 - Class Name: crane
Folder: n03179701 - Label: 10 - Class Name: desk
Folder: n02509815 - Label: 11 - Class Name: lesser panda, red panda, panda, bear cat, cat bear, Ailurus fulgens
Folder: n02730930 - Label: 12 - Class Name: apron
Folder: n03160309 - Label: 13 - Class Name: dam, dike, dyke
Folder: n02769748 - Label: 14 - Class Name: backpack, back pack, knapsack, packsack, rucksack, haversack
```

- First 400 from each folder selected for training
- Last 100 selected for testing
- For fine tuning I needed a validation set
- So split 400 training into 300/100 with the last 100 used for validation
- Final split per folder
- 300/100/100
- Final Split Overall
- 4500 Training /1500 Validation /1500

Handcrafted Features

Preprocessing Differences:

SIFT - Used histogram equalisation and masking to enhance image

ORB - Resized images to 256 x 256

| Model | Num_Words | Accuracy |
|----------------------|-----------|----------|
| SIFT Bag of Words | 50 | 21.67% |
| SIFT Bag of Words | 100 | 23.13% |
| SIFT Bag of Words | 150 | 21.47% |
| SIFT Bag of Words | 200 | 20.40% |
| ORB Bag of Words | 50 | 24.07% |
| ORB Bag of Words | 100 | 22.73% |
| ORB Bag of Words | 150 | 22.40% |
| ORB Bag of Words | 200 | 21.87% |

| Model | Num_Clusters | Accuracy |
|-----------------------|--------------|----------|
| SIFT Fisher Vector | 25 | 14.07% |
| SIFT Fisher Vector | 50 | 18% |
| SIFT Fisher Vector | 100 | 15.73% |
| ORB Fisher Vector | 25 | 20.60% |
| ORB Fisher Vector | 50 | 24.09% |
| ORB Fisher Vector | 100 | 22.67% |

Key Takeaways:

- SIFT was far more computationally expensive
- With BoW SIFT was more suited
- After 125 Key words overfitting happened
- The statistical power of GMM proved useful
- ORB and fisher vector outperformed other
- Still poor results with dependency on certain features and prone to default to guessing same label.

Linear Neural Network and Convolutional Neural Network Training Phase

Preprocessing:

- Dataset had mix of greyscale and RGB
- Converted all to RGB
- Batched into default size of 32 using DataLoader
- Transformed to tensor and normalised data

```
param_grid = {  
    'learning_rate': [0.001, 0.01],  
    'batch_size': [32, 64, 128, 256],  
    'dropout_rate': [0.3, 0.5],  
    'num_epochs': [40, 50, 60],  
    'optimizer': ['adam', 'sgd'] #
```

Training Process:

1. Start with a base model with simple number of layers.
2. Find downfalls and try improve on previous
3. Create a set of models and test
4. Choose a model to take for further improvement
5. Apply techniques like early stopping, hyperparameter tuning and batch normalisation

Techniques Used:

1. Early Stopping with Patience
2. Hyperparameter Tuning:
 - Optimiser Type
 - Learning Rate
 - Batch Size
 - Number of epochs
 - Dropout Rate
 - Learning Rate Scheduler
3. Data Augmentation
4. Batch Normalisation
5. Global Average Pooling (CNN)

```
Early stopping triggered after epoch 15  
Training with params: {'batch_size': 32, 'dropout_rate': 0.3, 'learning_rate': 0.001, 'num_epochs': 40, 'optimizer': 'sgd'}  
Epoch 1/40 - Train Accuracy: 14.56% | Train Loss: 2.5931 | Val Accuracy: 23.47% | Val Loss: 2.4200  
Epoch 2/40 - Train Accuracy: 23.33% | Train Loss: 2.4049 | Val Accuracy: 25.13% | Val Loss: 2.3331  
Epoch 3/40 - Train Accuracy: 26.47% | Train Loss: 2.3088 | Val Accuracy: 28.27% | Val Loss: 2.2701  
Epoch 4/40 - Train Accuracy: 29.58% | Train Loss: 2.2259 | Val Accuracy: 30.26% | Val Loss: 2.2208  
Epoch 5/40 - Train Accuracy: 31.40% | Train Loss: 2.1588 | Val Accuracy: 30.40% | Val Loss: 2.1934  
Epoch 6/40 - Train Accuracy: 33.47% | Train Loss: 2.0902 | Val Accuracy: 30.87% | Val Loss: 2.1894  
Epoch 7/40 - Train Accuracy: 35.46% | Train Loss: 2.0374 | Val Accuracy: 31.53% | Val Loss: 2.1789  
Epoch 8/40 - Train Accuracy: 38.02% | Train Loss: 1.9668 | Val Accuracy: 30.67% | Val Loss: 2.1668  
Epoch 9/40 - Train Accuracy: 38.40% | Train Loss: 1.9321 | Val Accuracy: 32.33% | Val Loss: 2.1523  
Epoch 10/40 - Train Accuracy: 40.40% | Train Loss: 1.8817 | Val Accuracy: 32.47% | Val Loss: 2.1548  
Epoch 11/40 - Train Accuracy: 42.84% | Train Loss: 1.8174 | Val Accuracy: 33.00% | Val Loss: 2.1569  
Epoch 12/40 - Train Accuracy: 45.71% | Train Loss: 1.7549 | Val Accuracy: 32.87% | Val Loss: 2.1523  
Epoch 13/40 - Train Accuracy: 45.10% | Train Loss: 1.7379 | Val Accuracy: 33.00% | Val Loss: 2.1751  
Early stopping triggered after epoch 14  
Training with params: {'batch_size': 32, 'dropout_rate': 0.3, 'learning_rate': 0.001, 'num_epochs': 50, 'optimizer': 'adam'}  
Epoch 1/50 - Train Accuracy: 11.76% | Train Loss: 3.1682 | Val Accuracy: 15.13% | Val Loss: 2.6058  
Epoch 2/50 - Train Accuracy: 15.71% | Train Loss: 2.6329 | Val Accuracy: 19.47% | Val Loss: 2.5227  
Epoch 3/50 - Train Accuracy: 17.40% | Train Loss: 2.5758 | Val Accuracy: 19.47% | Val Loss: 2.5353  
Epoch 4/50 - Train Accuracy: 19.27% | Train Loss: 2.5466 | Val Accuracy: 20.93% | Val Loss: 2.4841  
Epoch 5/50 - Train Accuracy: 20.56% | Train Loss: 2.5849 | Val Accuracy: 20.47% | Val Loss: 2.4777  
Epoch 6/50 - Train Accuracy: 21.36% | Train Loss: 2.4743 | Val Accuracy: 22.27% | Val Loss: 2.4530  
Epoch 7/50 - Train Accuracy: 20.73% | Train Loss: 2.4029 | Val Accuracy: 23.07% | Val Loss: 2.4311  
Epoch 8/50 - Train Accuracy: 21.89% | Train Loss: 2.4566 | Val Accuracy: 21.67% | Val Loss: 2.4551  
Epoch 9/50 - Train Accuracy: 22.29% | Train Loss: 2.4329 | Val Accuracy: 20.47% | Val Loss: 2.4413  
Epoch 10/50 - Train Accuracy: 21.80% | Train Loss: 2.4435 | Val Accuracy: 22.27% | Val Loss: 2.4655  
Epoch 11/50 - Train Accuracy: 22.62% | Train Loss: 2.4281 | Val Accuracy: 22.67% | Val Loss: 2.4278  
Epoch 12/50 - Train Accuracy: 22.27% | Train Loss: 2.4202 | Val Accuracy: 20.73% | Val Loss: 2.4654  
Epoch 13/50 - Train Accuracy: 22.36% | Train Loss: 2.4321 | Val Accuracy: 22.13% | Val Loss: 2.4325  
Epoch 14/50 - Train Accuracy: 22.62% | Train Loss: 2.4308 | Val Accuracy: 20.93% | Val Loss: 2.4315  
Epoch 15/50 - Train Accuracy: 24.40% | Train Loss: 2.3619 | Val Accuracy: 22.13% | Val Loss: 2.4030  
Epoch 16/50 - Train Accuracy: 23.33% | Train Loss: 2.3577 | Val Accuracy: 20.40% | Val Loss: 2.4700  
Epoch 17/50 - Train Accuracy: 25.11% | Train Loss: 2.3494 | Val Accuracy: 21.60% | Val Loss: 2.4007  
Epoch 18/50 - Train Accuracy: 23.82% | Train Loss: 2.3690 | Val Accuracy: 17.67% | Val Loss: 2.4846  
Epoch 19/50 - Train Accuracy: 24.78% | Train Loss: 2.3368 | Val Accuracy: 18.33% | Val Loss: 2.4929  
Early stopping triggered after epoch 20  
Training with params: {'batch_size': 32, 'dropout_rate': 0.3, 'learning_rate': 0.001, 'num_epochs': 50, 'optimizer': 'sgd'}  
Epoch 1/50 - Train Accuracy: 14.68% | Train Loss: 2.6038 | Val Accuracy: 23.93% | Val Loss: 2.4317  
Epoch 2/50 - Train Accuracy: 22.58% | Train Loss: 2.4098 | Val Accuracy: 27.67% | Val Loss: 2.3272  
Epoch 3/50 - Train Accuracy: 25.73% | Train Loss: 2.3208 | Val Accuracy: 29.93% | Val Loss: 2.2639  
Epoch 4/50 - Train Accuracy: 29.71% | Train Loss: 2.2232 | Val Accuracy: 30.40% | Val Loss: 2.2312  
Epoch 5/50 - Train Accuracy: 31.62% | Train Loss: 2.1635 | Val Accuracy: 30.67% | Val Loss: 2.2023
```

Linear Neural Network

Models:

LinearNNV1: Baseline

- Input Size: 12288 (Flattened Image)
- Hidden Layers: 2 → [512, 256]
- Activation: ReLU
- Dropout: 0.5 (after each layer)
- Output: 15 neurons (for 15 classes)

LinearNNV2: Adding Layers

- Hidden Layers: 4 → [1024, 512, 256, 128]
- Dropout: 0.3

LinearNNV3: Improved Feature Extraction

- Hidden Layers: 6 → [2048, 1024, 512, 256, 128, 64]
- Dropout Strategy: Gradually decreasing (0.5 → 0.2)
- Goal: Better generalisation with strong feature learning

Final Model

- V2
- Early Stopping
- SGD (0.001 LR)
- Batch Size = 128

| Model (architecture) | Epochs | Batch_Size | Optimiser | Data Augmentation | Early Stopping | Hyperparameter Tuning | Training Acc | Validation Acc | Test Acc | Test Loss |
|----------------------|--------|------------|------------------------------|-------------------|----------------|-----------------------|--------------|----------------|----------|-----------|
| V1 | 40 | 32 | Adam(lr=0.001) | No | No | No | 27.73% | 19.20% | 22.60% | 2.4754 |
| V2 | 40 | 32 | Adam(lr=0.001) | No | No | No | 30.18% | 18.07% | 18.20% | 2.6662 |
| V3 | 40 | 32 | Adam(lr=0.001) | No | No | No | 33.93% | 23.93% | 25.73% | 2.6344 |
| V1 | 50 | 64 | SGD(lr=0.001 Dropout=0.3) | No | Yes | Yes | 48.31% | 34.26% | 33.47% | 2.1243 |
| V2 | 60 | 128 | SGD(lr=0.01 Dropout=0.3) | No | Yes | Yes | 39.73% | 32.47% | 31.53% | 2.1900 |
| V3 | 40 | 32 | SGD(lr=0.01) Dropout=0.3 | No | Yes | Yes | 38.76% | 28.13% | 26.54% | 2.4587 |

| Model (architecture) | Epochs | Batch_Size | Optimiser | Data Augmentation | Early Stopping | Hyperparameter Tuning | Training Acc | Validation Acc | Test Acc | Test Loss |
|----------------------|--------|------------|---|-------------------|------------------|-----------------------|--------------|----------------|----------|-----------|
| V2 | 60 | 128 | SGD(0.001, momentum=0.9) Dropout Rate=0.3 | No | Yes (Patience=5) | Yes | 45.42% | 32.47% | 33.00% | 2.1522 |
| V2 | 60 | 128 | SGD(0.001, momentum=0.9) Dropout Rate=0.5 | No | Yes (Patience=5) | Yes | 32.09% | 29.73% | 30.00% | 2.1893 |
| V2 | 100 | 128 | SGD(0.001, momentum=0.9) Dropout Rate=0.3 | Yes | Yes (Patience=5) | Yes | 44.20% | 28.93% | 30.47% | 2.2116 |
| V2 | 100 | 128 | SGD(0.001, momentum=0.9) Dropout Rate=0.5 | Yes | Yes (Patience=5) | Yes | 35.00% | 27.33% | 27.67% | 2.2888 |

Convolutional Neural Network



Had a varying number of models ranging from V1-V8

V1-V5 focused on varying number of convolutional layers, dropout rates and fully connected layers

V6-V8 built on this but incorporated batch normalisation and global average pooling

Key Takeaways

- Early Stopping was effective
- Four Convolutional Layers was best amount
- Battle between overfitting and performance
- Adam Optimiser best with small learning rate
- Augmentation improved fit



Comparison of Best Models

| Model | Epochs/Number of Clusters | Architecture | Hyperparameters | Test Accuracy |
|----------------------------|---------------------------|--|--|---|
| ORB Fisher Vector | 50 | Fitted Gaussian Mixture Model Classified using LinearSVC() Output (15 classes) | Max_iter=1000, C=0.1 | 24.09% |
| LinearNNV2 | 60 | Input Size - 64 x 64 x 3 Flattening Layer - (12288) Four Hidden Layers = [1024, 512, 256, 128] ReLU Activation Dropout - 30% Output(15 classes) | Batch_size = 128 Optimiser = SGD Learning Rate=0.001 Momentum = 0.9 Loss = CrossEntropy Dropout Rate=0.5 Early Stopping Patience = 5 | 33.00% |
| CNNV7 | 100 | Input (64×64×3) → Conv1 (16 filters, 3×3, ReLU, BN, MaxPool 2×2) → Conv2 (32 filters, 3×3, ReLU, BN, MaxPool 2×2) → Conv3 (64 filters, 3×3, ReLU, BN, MaxPool 2×2) → Conv4 (128 filters, 3×3, ReLU, BN, MaxPool 2×2) → Global Avg Pool → Fully Connected Layer (128 → 15) → Dropout (20%) → Output (15 classes). BN (Batch Normalisation) | Batch_size= 64 Optimiser = Adam Learning Rate=0.001 Scheduler = StepLR(10,0.1) Loss=CrossEntropy Dropout Rate = 0.2 Early Stopping Patience=3 Data Augmentation=No | 62.27% |
| CNNV7 Data Augmentation | 100 | Same as above | Same as above except Data Augmentation=True | 58.67% (Neater Train/Val loss plot) |

Error Analysis

Linear NN vs Handcrafted Features

- Handcrafted Model: Relied on dominant features → High recall, low precision (e.g., guinea pigs & spiny lobsters).
- Linear NN: Learned more balanced features → Better precision & recall, improved performance on labels like 2 & 13.

Linear NN vs CNNs

- Linear NN Mistakes: Predicted chimpanzee (4) as backpack (14) → Features not informative.
- CNN Advantage: Avoided major mistakes, instead misclassifying boa constrictor (1) as spiny lobster (7) → More logical errors based on shared visual features.

Key Test Case: Bucket & Pail (Label 3)

- Worst-performing category across all models due to distractions in images.
- Handcrafted Model: Focused on irrelevant features (e.g., eyes).
- CNN: Extracted useful spatial features, leading to far more correct predictions.

